

Enterprise MCP Service Bus

A reference architecture for enterprise agentic access governed by Client-Bound Entitlements

Document class	Technical-executive architecture paper
Domain	Enterprise integration, identity, and authorization for agentic systems over the Model Context Protocol (MCP)
Nature	Reference architecture + executable reference implementation, maintained in this repository
Audience	Security and integration architects, platform engineering, technical and executive leadership
Repository	github.com/mcloh/Enterprise-MCP-Service-Bus

Contents

- 0. Executive summary
- 1. The problem: AI agents as a new class of enterprise consumer
- 2. Core idea and structural analogy
- 3. Core mechanisms (the inventive core of the architecture)
- 4. Component architecture
- 5. End-to-end operational flows
- 6. Threat model and defense in depth
- 7. Implementation evidence: from architecture to a working system
- 8. Comparative positioning
- 9. Formal set of invariants
- 10. Maturity model and adoption path
- 11. Glossary
- 12. Final synthesis

0. Executive summary

Agentic AI systems — agents that plan, decide, and call tools autonomously — are moving from isolated prototypes to regular consumers of enterprise systems: CRMs, ERPs, payment platforms, customer data, HR systems. The Model Context Protocol (MCP) has emerged as the de facto standard for connecting these agents to "tools" — functions that a language model can discover and invoke.

This new consumption pattern creates a problem that traditional enterprise integration never had to solve: **the API consumer is now a probabilistic component**. An agent can be manipulated by malicious content embedded in a document, an email, a web page, or the output of a previous tool call — an attack known as *prompt injection*. If authorization depends, even indirectly, on any information the agent itself processes or declares, a successful attack against the model translates directly into a privilege escalation against the business.

This document describes an architecture — the **Enterprise MCP Service Bus** — that solves this problem by moving the trust boundary outside the agent. The central idea, from which everything else derives, is simple to state and hard to violate by accident:

The authenticated identity of the MCP Client — never the content processed by the agent — defines the maximum ceiling of capabilities a session can exercise. No information controlled by the agent can widen that ceiling; it can only narrow it.

From this single principle, the architecture builds an enterprise bus — inspired by the historical role of Enterprise Service Buses (ESBs), but designed for semantic consumption by agents — that:

- publishes enterprise capabilities once and reuses them through governed subscription, instead of every team building its own ad hoc connector;
- delivers each agent a **personalized catalog**, not the company's entire catalog, simultaneously reducing risk, model context cost, and tool-selection error rate;
- re-authorizes **every** execution call independently of the listing, closing the classic gap between "what the agent can see" and "what the agent can do";
- allows personalization, recommendation, and journey orchestration (including computing a *Next Best Action*) **without ever letting relevance turn into authorization**;

- produces an end-to-end correlated audit trail, from the business event to the policy decision to the backend call.

The practical result is that a cognitive compromise of the agent — the failure mode the entire industry already assumes is inevitable — stops being an event of unbounded privilege escalation and becomes, at worst, **abuse contained inside a previously sized and auditable compartment**.

The architecture is not merely conceptual: the reference implementation maintained in this repository materializes every mechanism described below as executable code, validated against real dependencies (a policy engine, an identity provider, a state-graph orchestrator, a real language model) and exercised by a suite of more than 260 automated tests, including adversarial tests that physically reproduce the attack scenarios discussed in Section 6. Section 7 details this evidence.

1. The problem: AI agents as a new class of enterprise consumer

In a mid-size to large organization, business capabilities are already fragmented across dozens of systems: REST and SOAP APIs, serverless functions, ERPs, CRMs, databases, legacy systems, and third-party SaaS. Historically, every new application that needed these systems went through an API gateway, an ESB, or a service mesh — layers that solved routing, transformation, and security problems for **deterministic** consumers: code that calls exactly the API it was programmed to call.

AI agents break that premise in two simultaneous ways:

1. **The consumer is no longer deterministic.** An agent decides, at run time, which tools to invoke and with which arguments, from a context that may include untrusted third-party text.
2. **The tool catalog becomes part of the agent's behavior.** Unlike an undocumented HTTP endpoint, a tool described in `tools/list` is typically folded into the model's context and directly influences what it will attempt to do. Controlling what appears in that listing stops being merely a documentation concern and becomes a security decision, a context-cost decision, and a tool-selection-quality decision, all at once.

Without a dedicated enterprise layer, the pattern is the same in practically every organization that adopts agents in a decentralized way: each team implements its own MCP Server, with its own authentication model, exposes catalogs larger than necessary "for convenience," grants backend credentials directly to the agent, and discovers too late that revoking or auditing that access is hard. The diagram below contrasts that scenario with the result of introducing a single enterprise bus as the point of publication, discovery, and — above all — policy enforcement.




Figure 1 — Without a governance layer, every agent accumulates its own integrations and credentials ($N \times M$ connections, each a distinct risk and audit vector). With the bus, all agentic traffic passes through a single policy-enforcement point before reaching any enterprise system.

The problem, then, is not "how to connect an LLM to a tool" — MCP already solves that at the protocol level. The problem is **how to make hundreds or thousands of enterprise capabilities available to many agentic consumers consistently, in a governed and secure way**, accepting as a design premise that the consumer may be cognitively compromised.

2. Core idea and structural analogy

The architecture deliberately assumes the historical role of an Enterprise Service Bus, adapted to a semantic, probabilistic consumer. The difference is not one of layers — routing, mediation, transformation, and service resolution are all still present — but of **where the trust boundary sits** and of **what flows between agent and bus**: not just data, but capability descriptions that the agent itself uses to decide what to do next.





Figure 2 — The traditional ESB and the Enterprise MCP Service Bus share the same skeleton (gateway → bus → systems). The structural difference is that the MCP Client becomes an explicit identity component between the agent and the gateway, and what travels left to right is not a fixed payload but a semantic capability description that the consumer itself reasons over.

That difference — capabilities described semantically and consumed by a probabilistic planner — is why this architecture cannot simply inherit the trust model of a traditional API Gateway. It must treat **what the agent can discover** and **what the agent can execute** as two distinct controls, and it must formally guarantee that nothing inside the agent's session — prompt, memory, prior tool output, retrieved content — can alter the authorization ceiling established outside it. Those two points are the subject of the next section.

The consolidated view below shows how the security core (the bus) relates to four additional operational responsibilities that the architecture organizes around it: capability-offer governance, governed personalization, journey orchestration, and omnichannel execution.




Figure 3 — Consolidated view. Two chains run through the diagram: a supply chain (service owner → pipeline → registry → fabric) and a decision chain (profile + entitlement + filtered offerings → decision context → NBA → execution). They meet only at the bus's governed MCP contract; neither one has authority to bypass it.

From here on, the color legend is stable for the rest of the document: **amber** marks the deterministic security boundary (authentication, PEP, PDP); **teal** marks the bus's trusted control plane (fabric, routing, resolution); **violet** marks capability governance and lifecycle; **terracotta** marks the untrusted/probabilistic zone (agent, external content, channels); **slate blue** marks enterprise systems of record.

3. Core mechanisms (the inventive core of the architecture)

This section describes, one by one, the mechanisms that make the architecture technically different from "putting a proxy in front of an MCP Server." Each mechanism is presented with the problem it solves, its precise operation, and the formal invariant it guarantees.

3.1 Client-Bound Maximum Entitlement

Problem. In naive implementations, authorization is evaluated from information the agent itself, the input payload, or the system prompt declare — for example, an `agent_id` or `role` field read from the request body. Any information on that path is, by definition, under the influence of content the model processes, and is therefore manipulable via *prompt injection*.

Mechanism. Every MCP Client holds an identity **authenticated by a cryptographic mechanism independent of the conversation's content** (OAuth 2.1 client credentials, workload identity, mTLS, or equivalent). That identity — never the self-declared `clientInfo`, never a payload field — is the only key used to resolve the `MaximumEntitlement`: the maximum set of capabilities that session can, under any circumstance, ever execute. Every additional restriction (execution context, end-user attributes, risk policy) can **shrink** that set; none can widen it.




Figure 4 — The authorization ceiling originates exclusively from the client's authenticated identity; every other input can only narrow the effective set, never widen it.

Invariant. `EffectiveCapabilities(client, context) ⊆ MaximumEntitlement(client)`, for any context, including a context under attack. Formally, the effective set is the intersection of the client's ceiling with user, runtime, risk, and resource constraints — never a union.

Technical effect. A cognitive compromise of the agent (direct or indirect prompt injection, RAG poisoning, manipulated tool output) can no longer produce a capability outside that client's pre-approved compartment, no matter how convincing the malicious instruction is.

3.2 Separation between discovery authorization and execution authorization

Problem. Hiding a tool from a listing (`tools/list`) reduces exposure, but does not by itself prevent a consumer who already knows the tool's name from invoking it directly via `tools/call`. Treating "not listed" as synonymous with "protected" is a common and exploitable mistake.

Mechanism. The architecture defines two independent controls, applied at different moments in the call lifecycle, each re-evaluated from the same authenticated identity:




Figure 5 — Discovery authorization and execution authorization never share the same implicit verdict: a tool hidden from the catalog is still re-evaluated from scratch if anyone attempts to call it directly.

Invariant. `Protected = Hidden from unauthorized discovery AND Denied when unauthorized execution is attempted`. Neither control, on its own, satisfies the definition.

Technical effect. Eliminates the attack class in which a consumer obtains the name of a sensitive tool — through reverse engineering, documentation leakage, or simple trial and error — and invokes it without ever having gone through the filtered listing.

3.3 Decision Context as a governed personalization envelope

Problem. Personalization and journey orchestration (what to recommend, to whom, on which channel, right now) traditionally live in systems separate from the authorization layer, which creates a recurring temptation: let a propensity model, a profile score, or a business rule decide directly what the agent may do — mixing relevance with permission.

Mechanism. All the information needed for an orchestration decision is composed into a single, versioned envelope — the `DecisionContext` — **before** it reaches the component that decides the next action. That envelope combines a profile view (replaceable, produced by any segmentation engine), the authenticated identity, the already-resolved entitlement limits, the set of offerings already filtered by the authorization ceiling, and the runtime context (channel, consent, moment). The essential point is the order: the entitlement filter happens **before** the envelope is exposed to the orchestrator, never after.




Figure 6 — The Decision Context is assembled from components already restricted to the entitlement ceiling; the orchestrator never receives, and therefore can never widen, a universe larger than that.

Technical effect. Swapping the personalization engine (from deterministic rules to a ranking model, from one clustering algorithm to another) requires no change to the authorization core, because the orchestrator's input contract is already, by construction, a subset of what is permitted. Personalization becomes a pluggable component without becoming a privilege-escalation surface.

3.4 Monotonic restriction: relevance never widens authorization

Problem. Recommendation, clustering, and *decisioning* systems are, by nature, dynamic and probabilistic — exactly the kind of component this architecture treats as untrusted for authorization purposes (Section 3.1). An explicit mechanism is needed to prevent those components from indirectly introducing a capability outside the client's ceiling.

Mechanism. The architecture defines four sets with distinct authorities and a strictly decreasing containment relation between them. No later stage can reintroduce an element removed by an earlier stage.




Figure 7 — Strict containment across the four sets the architecture never lets collapse into one. Each layer can only shrink the one before it; the authority deciding each layer is different (registry, entitlement manager/PDP, offering filter, orchestrator).

Set	Question it answers	Authority	Can it widen the one outside it?
GlobalCatalog	What has been published and is active?	Registry + Fabric	—
MaximumEntitlement	What is this client's ceiling?	Entitlement Manager + PDP	No
FilteredOfferings	What is eligible right now?	Offering Filter	No
RankedCandidates / NBA	What is most relevant right now?	Service Orchestrator	No

Invariant. $\text{RankedCandidates} \subseteq \text{FilteredOfferings} \subseteq \text{MaximumEntitlement} \subseteq \text{GlobalCatalog}$. A propensity score, a profile-cluster change, or a new journey rule may reorder or shrink candidates; none of them is, by architectural definition, a *security principal* capable of creating entitlement.

3.5 Cache scoped to the authorization context

Problem. Tool listings are expensive to recompute and naturally attractive to cache. But an entitlement-filtered catalog is, by definition, a response that is **private to that authorization context**; a cache naively shared across different clients leaks exactly the information discovery filtering was meant to protect.

Mechanism. Every entitlement-bound listing response is treated as non-shareable across distinct authorization contexts. The cache key incorporates the client's authenticated identity, the authorization context, and the current policy version; no entry is ever served to an identity different from the one it was computed for, and policy changes invalidate the affected entries.




Figure 8 — Two different identities never share a single catalog cache entry, even when the underlying call is otherwise identical.

Technical effect. Eliminates a class of "performance-accident" capability leak, in which a caching optimization silently reintroduces the exact information leak that discovery filtering was designed to prevent.

3.6 Separation between inbound identity and outbound identity

Problem. It is tempting to propagate, unchanged, the token a client used to authenticate to the bus as the credential for calling a backend system ("token passthrough"). This blends trust domains (*audience confusion*) and creates a classic *confused deputy* pattern: the backend ends up implicitly trusting whatever the gateway forwards.

Mechanism. The credential used to enter the bus and the credential used to reach a backend system are **distinct** relations, mediated by a dedicated identity exchange. The client obtains a token with audience restricted to the bus; once the gateway authorizes the call, the integration layer separately obtains a credential appropriate to that specific backend — never the agent's original token.




Figure 9 — The inbound credential never crosses, unmodified, the boundary into the backend. Every trust hop carries its own credential, with its own audience.

Technical effect. A stolen agent credential does not automatically translate into a backend credential; the blast radius of a leak stays bounded to the boundary where the leak occurred, not the entire chain of systems behind it.

3.7 Governed capability publication (the capability supply chain)

Problem. Capabilities registered manually, without a defined owner, without risk classification, and without a policy gate, tend to accumulate as invisible technical debt: nobody is quite sure what is exposed, to whom, or under what risk.

Mechanism. Every capability is born from a declarative manifest (canonical name, domain, version, input/output schema, owner, risk and data classification, side effects, idempotency, required entitlements, approval requirements) and only becomes resolvable by the bus after passing through a publishing pipeline with mandatory schema, risk, and policy gates.




Figure 10 — No capability reaches the "active and resolvable" state without mechanically passing through all four gates. The reference implementation in this repository enforces this rule at the code level: nothing reaches the active state without first passing review.

Technical effect. Turns the surface of capabilities exposed to the agentic ecosystem into a managed asset with lifecycle, versioning, and an accountable owner — not a list that grows organically through manual registration.

3.8 End-to-end correlated traceability

Problem. In a system with multiple decision components (profile, entitlement, offering filter, orchestrator, gateway, backend), the absence of a stable correlation key makes it practically impossible to reconstruct, after the fact, why a specific action was taken or denied.

Mechanism. Every relevant decision — profile resolution, entitlement resolution, policy decision, NBA computation, MCP call, backend result — carries versioned, correlatable identifiers (profile version, entitlement version, decision id, policy-decision id, MCP request id) propagated as first-class fields through the entire chain, never reconstructed *after the fact* by heuristic timestamp correlation.

Technical effect. Every authorization decision is deterministic and auditable (Invariant 10, Section 9): given a business event, it is possible to reconstruct exactly which policy version, which entitlement, and which orchestration decision led to a specific outcome — an indispensable requirement both for incident investigation and for demonstrating compliance.

4. Component architecture

The table and diagram below consolidate each component's responsibility and the trust boundaries between them. Three logical zones organize the system: the untrusted/probabilistic zone (the agent and everything it processes), the deterministic security zone (authentication, PEP, PDP, audit), and the trusted enterprise integration zone (fabric and backend systems).




Figure 11 — The trust boundary sits between the MCP Client and workload authentication — never inside the zone where the agent reasons. No arrow crosses that boundary in reverse.

Component	Core responsibility	Authority over
Agent / Agent Host	Interpret intent, plan, select tools, orchestrate the model session	No security decision
MCP Client	Authenticate to the bus; carry the identity that defines the privilege ceiling	The maximum compartment for that session
MCP Gateway / PEP	Enforce authentication, discovery filtering, execution re-authorization, rate limiting, audit	Enforcement — never delegates the decision to the model
PDP / Policy Engine	Evaluate entitlement policy (does the tool belong to the ceiling?) and transaction policy (is the argument acceptable?)	Deterministic ALLOW / DENY / REQUIRE_APPROVAL decision
Global Capability Registry	Complete administrative catalog, with lifecycle, owner, and risk classification	Definition of what exists and is active
Enterprise MCP Fabric	Routing, protocol mediation, transformation, backend resolution	Execution of an already-authorized capability
Entitlement Manager	Resolve, version, and revoke <code>MaximumEntitlement</code> by authenticated identity	The maximum menu per client

Component	Core responsibility	Authority over
Profile Intelligence	Produce a replaceable profile view (segments, scores, attributes)	Relevance — never authorization
Offering Filter	Intersect active catalog, entitlement, and context to produce eligible candidates	The universe the NBA may choose from
Service Orchestrator	Maintain journey state and compute the Next Best Action	The intended action — not authorization to execute it
Agent Runtime	Dispatch the NBA, manage session, channel, and handoff between agents	Coordinated execution, never credential issuance

The operating principle that ties the whole table together: **each component has authority over exactly one question**. The Orchestrator decides *which* action to attempt; the Gateway/PDP decides *whether* it may be executed; the backend keeps deciding over its own resources. None of these three roles is interchangeable, and a compromise of any one of them does not automatically grant the authority of the other two.

5. End-to-end operational flows

5.1 Filtered discovery (tools/list)




Figure 12 — Identity is authenticated before any filtering; the agent never chooses, declares, or influences the entitlement that bounds what it will see.

5.2 Re-authorized execution (tools/call)




Figure 13 — `tools/call` is re-authorized from scratch, in two steps (entitlement, then transaction), independently of what `tools/list` returned.

5.3 Orchestrated flow — from profile to execution, without merging authorities




Figure 14 — Profile, entitlement, and filtered offerings feed the NBA computation, but the orchestrator's decision keeps being re-authorized, call by call, by the same Gateway/PDP pair as the simple execution flow.

6. Threat model and defense in depth

The architecture does not assume that *guardrail* techniques will eliminate *prompt injection* — it assumes the opposite: that a cognitive compromise of the agent is a plausible, recurring failure mode, and asks only: **if the agent is compromised, what is the most it can do?** The desired answer is: at most what the MCP Client was already authorized to do, still subject to transaction policy, human approval where applicable, and the backend's final authority.

6.1 Scenario — prompt injection attempting privilege escalation




Figure 15 — The PEP never consults the model to decide; the denial happens solely because the requested capability does not belong to the client's ceiling, no matter how convincing the malicious instruction was.

The architecture distinguishes two classes of residual risk, with different mitigations for each:

Class	Definition	Example	Primary mitigation
Privilege escalation	Attempting to access a capability outside the entitlement	<code>Sales.Read</code> attempting <code>finance.createPayment</code>	PEP + client entitlement (Section 3.1)
Privilege abuse	Maliciously using an already-legitimate capability	<code>email.send</code> , authorized, instructed to leak data to an external recipient	Argument/resource policy, destination allowlist, human approval, transaction limits

Privilege escalation is resolved structurally by the entitlement model; *privilege abuse* requires a second line of defense over arguments and resources, because the tool itself is legitimate — the problem lies in its use.

6.2 Defense in depth




Figure 16 — No single layer resolves the whole problem; the guarantee comes from composition. A failure at L5 (argument policy), for instance, is still contained by L6 (approval) and L7 (backend authorization).

6.3 Blast-radius containment

The design decision with the largest practical effect on residual risk is the **size of the compartment per client**. A single, very-high-privilege client shared by many agents concentrates risk; multiple clients segmented by domain and risk level bound the impact of any individual compromise to a predictable radius.




Figure 17 — The same number of agents produces radically different blast radii depending solely on how MCP Clients are segmented. Segmenting by domain and risk classification — not by agent — is the design variable with the largest effect on the expected impact of a compromise.

Every published capability carries an explicit risk classification, used both to segment clients and to decide where human approval is mandatory:

Tier	Nature	Example	Expected control
R0	Discovery / public	non-sensitive metadata	authentication optional depending on context
R1	Read	search, lookup	entitlement + audit
R2	Write	create/update a record	entitlement + argument policy

Tier	Nature	Example	Expected control
R3	High impact	approve credit, execute a payment	dedicated entitlement + policy + approval
R4	Destructive / restricted	terminate an account, delete a record	dedicated client + strong authentication + approval + backend controls

7. Implementation evidence: from architecture to a working system

An architecture document carries real technical weight when every mechanism described in the previous sections corresponds to a component that exists, runs, and is exercised by tests against real dependencies — not mocks. The reference implementation maintained in this repository fully covers the mechanisms of Sections 3 through 6, with the following composition:

7.1 Real dependencies exercised

Dependency	What it validates	Nature of the validation
Policy engine (real execution, not simulated)	The two decision policies described in Section 3.1 — entitlement and transaction	Policy unit tests + end-to-end integration tests
OAuth 2.1 / OIDC identity provider (containerized, real)	Client credential issuance and validation, <i>token exchange</i> for outbound identity (Section 3.6)	Integration tests against the real provider, including negative audience scenarios
State-graph orchestrator with persistent checkpointing	The profile → entitlement → filtered offerings → NBA → approval flow (Section 5.3), including pausing and resuming approvals	Unit and end-to-end tests, including proof that resuming a paused checkpoint does not re-execute already-completed nodes
Real language model, with tool-calling	The prompt-injection containment scenario (Section 6.1), physically reproduced against a real, not simulated, model	End-to-end test with a real tool call
Full containerized environment	Network segmentation between domains, fail-closed behavior, correlated tracing across distinct services	Full stack bring-up and verification that no domain exposes a port beyond the gateway
Self-hosted LLM observability platform	Real export of agent execution telemetry (tool-call and generation spans)	Functional validation against the real stack

7.2 Extent of test coverage

The reference implementation's test suite exceeds **260 automated cases**, covering four categories: unit tests for each component, contract tests between components, end-to-end tests against real dependencies, and a dedicated **adversarial testing** category that physically reproduces the attack scenarios described in Section 6 — including direct privilege-escalation attempts via *prompt injection*, direct calls to unlisted tools, *replay* of already-consumed approvals under real multi-thread concurrency, and formal verification (via thousands of generated cases) that the offering filter never reintroduces a capability outside the client's entitlement.

A dedicated subset of **security acceptance tests**, run against the full containerized environment, verifies end to end the architecture's core guarantees: unauthorized discovery never returns a capability outside the ceiling; a direct, unauthorized call is denied even when the exact tool name is known; no prompt content alters the maximum entitlement ceiling; policy revocation blocks further executions; direct access to the *fabric* bypassing the gateway fails; and simulating a stolen client credential demonstrates that the attacker never exceeds that specific client's entitlement.

7.3 Example domains and risk coverage

The reference implementation demonstrates segmentation by domain and by risk (Section 6.3) with two complete business domains — a sales domain and a finance domain — covering capabilities across the first four risk levels (R1 through R4), including an R3 capability with a complete human-approval cycle: initial denial, external out-of-band approval, execution, and a correctly denied identical retry via *replay* (the same approval is never consumed twice).

7.4 Engineering findings, reported honestly

A serious reference implementation also documents what was **not** solved, rather than presenting only the happy path. Two real concurrency findings, identified and fixed during implementation, illustrate the kind of detail that only surfaces once an architecture is turned into code:

- A real race condition in the high-risk approval gate: under concurrent simultaneous attempts, two requests could both observe the same approval as "granted" before either marked it as "consumed" — in theory allowing a double use of a single-use approval. Fixed with a dedicated critical section and proven under real concurrent load.
- An audit-sink failure could, before the fix, propagate as a server error on an operation that had already been decided — violating the principle that an observability failure must not block a low-risk decision that has already been made.

Structural limitations are also stated explicitly: the reference implementation does not include a dedicated request anti-*replay* mechanism (mitigated by short-lived credentials and complete auditing) nor rate limiting at the gateway (mitigated only by after-the-fact detectability, never by real-time prevention) — both described as necessary production extensions, not hidden gaps.

8. Comparative positioning

Dimension	Traditional ESB	API Gateway	Enterprise MCP Service Bus
Primary consumer	Applications	Applications / APIs	Agents / MCP Clients
Nature of the contract	Service / message	HTTP API	Semantic capability (tool)
Discovery	Service catalog	API catalog	Agent-visible catalog, filtered by entitlement
Semantic description	Limited	OpenAPI / metadata	Central — directly drives the consumer's planning
Policy enforcement	Common	Central	Mandatory, in two stages (discovery and execution)
Prompt injection as a threat	Not applicable	Indirect	Central design threat model
Catalog filtering by entitlement	Uncommon	Possible	Mandatory architectural element
Blast radius per identity	Possible	Possible	Central design principle (Section 6.3)

The difference the table alone does not capture is qualitative: an ESB or an API Gateway do not need to assume their consumer can be fooled by text. This architecture assumes exactly that as a design premise — and that premise is what determines each of the eight mechanisms in Section 3.

9. Formal set of invariants

The rules below are the architecture's non-negotiable contract. Any implementation, extension, or integration that violates one of them no longer satisfies this architecture, no matter how sophisticated the intelligence layer built on top of it is.

#	Invariant
1	The agent is not a security authority.
2	The authenticated identity of the MCP Client defines the maximum set of capabilities.
3	Information controlled by the agent can restrict, but never widen, authorization.
4	<code>EffectiveCapabilities</code> $\subseteq$ <code>ClientEntitlement</code> , always.
5	Discovery filtering is minimization; execution authorization is the real enforcement boundary.
6	Authorization of a tool does not imply unrestricted authorization of the transaction it performs.
7	The Gateway must not be bypassable by any network path or protocol.
8	The backend's trust in the bus stays bounded and explicit — never blind passthrough.
9	A compromised client must have a predictable, sized blast radius.
10	Every security decision is deterministic and auditable.
11	Profile intelligence may rank or reduce candidates; it may never grant capabilities.
12	A Next Best Action is an orchestration intent, not an authorization decision.
13	Capabilities are published once and reused only through governed subscription and enforcement.

10. Maturity model and adoption path

The architecture is designed for incremental adoption — no stage requires the next one to already exist.




Figure 18 — Each phase delivers standalone value. An organization can operate stably at Phase 1 or 2 indefinitely before moving forward.

Phase	Core deliverable
0 — Principles and contracts	Security invariants, capability metadata, client-profile model, identity model, audit schema
1 — Read-only pilot	One domain, a handful of R1 capabilities, filtered discovery, execution re-authorization, full audit
2 — Write operations	Argument policy, backend authorization, transaction limits, idempotency, selective approval
3 — Multi-domain fabric	Domain servers, central registry, namespace governance, lifecycle automation
4 — Governed personalization and orchestration	Versioned profile contract, Entitlement Manager, Offering Filter, Decision Context, Service Orchestrator, Agent Runtime, end-to-end correlation
5 — Enterprise scale	Automated client provisioning, policy as code, entitlement review, advanced anomaly detection, cost governance

11. Glossary

Agent — AI component that interprets context, plans, and requests execution of capabilities. Not a security authority.

Agent Runtime — Layer that dispatches the NBA and manages session, minimal context, and handoff between agents and channels.

Blast radius — Maximum expected impact if an identity or component is compromised.

Capability / Tool — A semantic operation exposed to an MCP consumer.

Client-Bound Capability View — MCP catalog filtered according to the authenticated client's entitlement.

Decision Context — Versioned envelope carrying profile view, authenticated identity, entitlement limits, filtered offerings, and runtime context, delivered to the Service Orchestrator.

Effective Capabilities — Subset of the Maximum Entitlement available in a specific context.

Entitlement Manager — Component that resolves, versions, and revokes the Maximum Entitlement and the personalized MCP menus derived from it.

Enterprise MCP Fabric — Integration layer that publishes, resolves, and routes MCP capabilities over enterprise services.

Global Capability Registry — Complete administrative catalog of governed capabilities, with lifecycle and ownership.

Maximum Entitlement — Capability ceiling associated with an authenticated MCP Client.

Next Best Action (NBA) — The action selected by the Orchestrator for a given journey context and objective; it does not, by itself, constitute execution authorization.

Offering Filter — Component that intersects the active catalog, entitlement, and context to produce eligible capabilities.

PDP — Policy Decision Point — Component that computes authorization decisions from policy and trusted attributes.

PEP — Policy Enforcement Point — Component that deterministically enforces authorization decisions.

Privilege abuse — Malicious or improper use of a legitimately granted capability.

Privilege escalation — Acquisition of a capability outside the entitlement.

Profile Intelligence — Replaceable capability that derives profile segments, clusters, or scores; informs relevance, never grants authorization.

Prompt injection — Attack that introduces malicious instructions into the context processed by a language model.

Service Orchestrator — Component that coordinates journey state and objective and computes the Next Best Action from Decision Context, rules, guardrails, and memory.

12. Final synthesis

The shift in perspective this architecture proposes can be summarized in a single sentence: **the agent does not receive access to an enterprise bus and then decide what to do with it; the MCP Client receives, in advance, an already-governed compartment of capabilities, and the agent can only operate within that compartment.**

That inversion — moving the authorization decision before and outside the layer where the model reasons — is what makes it possible to use MCP as a shared enterprise interface, accessible to many agents and use cases, without transferring the organization's root of trust to a probabilistic component. Around that core, the architecture organizes governed capability publication, replaceable personalization, journey orchestration, and omnichannel execution — each of these layers pluggable and evolvable, none of them with authority to change the authorization ceiling established by the client's identity.

The ultimate goal is not to make the agent trustworthy. It is to make it safe that it is not — and the reference implementation maintained in this repository exists precisely to demonstrate that this guarantee can be built, tested, and operated as real software, not merely described as an architectural principle.
